

Best Practices for Container Based Design

Microservices and Kubernetes-based Principles and Patterns






Eric Deandrea

Senior Principal Developer Advocate
edeandrea@redhat.com

Agenda

- Container Design Principles
- Kubernetes Patterns
 - ✧ Foundational Patterns
 - ✧ Structural Patterns
 - ✧ Configurational Patterns
 - ✧ Behavioral Patterns
 - ✧ Advanced Patterns





<https://k8spatterns.io>

<https://github.com/k8spatterns/examples>

<https://www.redhat.com/en/engage/kubernetes-containers-architecture-s-201910240918>


<https://developers.redhat.com/blog/2020/06/08/commit-to-excellence-java-in-containers>



Container Design Principles


Single Concern Principle (SCP)

- Every container should address a single concern and do it well
- Use sidecars and/or init-containers to combine multiple containers into a single pod
 - ✱ Each container still handles a single concern




Process Disposability Principle (PDP)

- A running container should be as ephemeral as possible and ready to be replaced by another container instance at any time





Runtime Confinement Principle (RCP)

- Containers have multiple dimensions at runtime, such as memory and CPU



Does Image Size Really Matter?



1 x 980 MB base image



10 x 20 MB project specific image


$$980 + 10 * 20 = 980 + 200 = \mathbf{1180\ MB}$$

≈ **200 MB**



- Total image size does not matter!
- Base image size does not matter!
- What really matters is the size of frequently changing layers!

Construct Application Layers the Right Way



Construct Application Layers the Right Way

980 MB layer

FROM registry.redhat.io/ubi8/openjdk-17:latest

10 MB layer

COPY target/dependencies /app/dependencies

9 MB layer

COPY target/resources /app/resources

1 MB layer

COPY target/classes /app/classes

ENTRYPOINT java -cp /app/dependencies/*:/app/resources:/app/classes my.app.Main



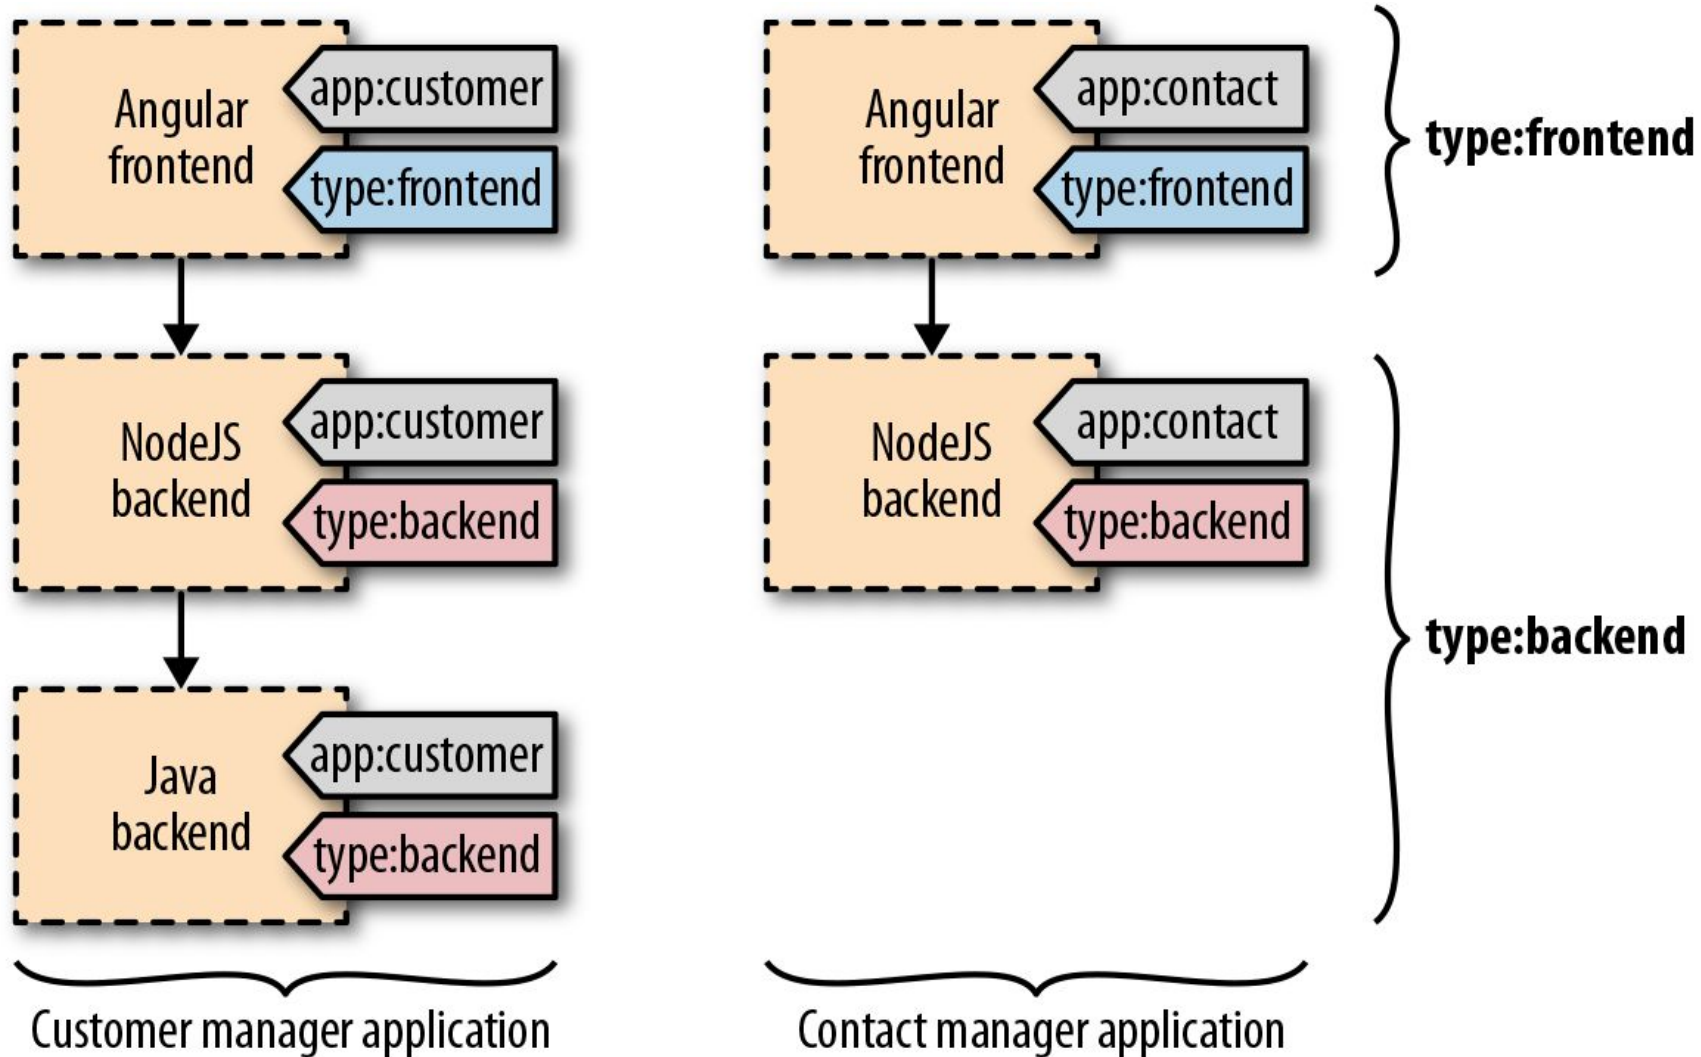
Foundational Patterns

Automatable Unit

How can we create and manage applications with Kubernetes.

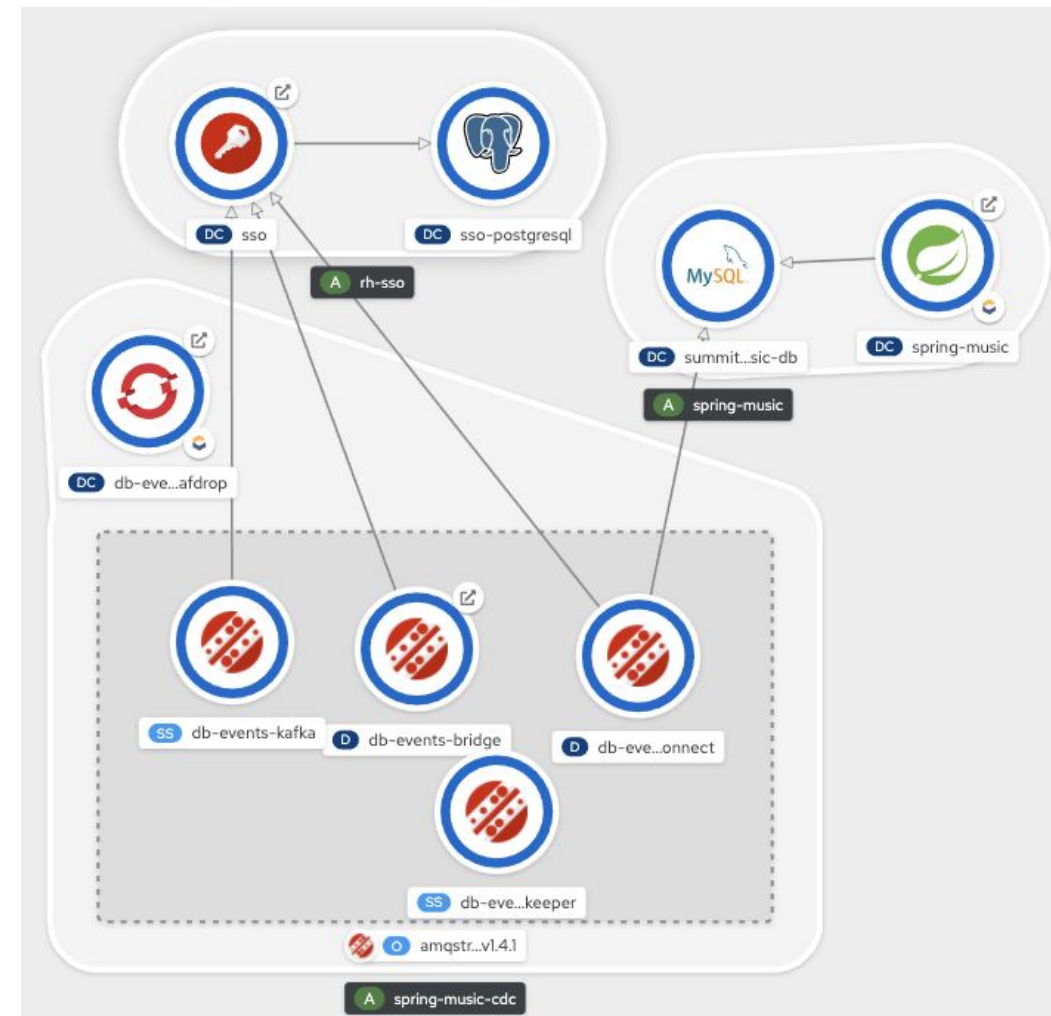
- Pods
 - ✱ Atomic unit within Kubernetes
- Services
 - ✱ Entry point to pods
- Grouping via Labels, Annotations, and Namespaces

Labels



Special Labels/Annotations

- app.kubernetes.io/part-of
- app.openshift.io/runtime
- app.openshift.io/runtime-version
- app.openshift.io/vcs-url
- app.openshift.io/vcs-ref
- app.openshift.io/connects-to



Predictable Demands

How can we handle resource requirements deterministically?

- Declared requirements
 - Scheduling decisions
 - Capacity planning
 - Matching infrastructure services
- Runtime dependencies
 - Persistent Volumes
 - Host ports
 - Dependencies on **ConfigMaps** and **Secrets**

Resource Profile Requests & Limits

- **Resources**
 - CPU, Network (compressible)
 - Memory (incompressible)
- **App**
 - Declaration of resource requests and limits
- **Platform**
 - Resource quotas and limit ranges

Quality-of-Service Classes

- **Best Effort**
 - No requests or limits
 - Lowest priority
- **Burstable**
 - requests < limits
- **Guaranteed**
 - requests == limits
 - Highest priority

Health Probe

How to communicate an application's health state to Kubernetes?

- **Process health checks**
 - Checks running process
 - Restarts container if container process died
- **Application provided Health Probes**
 - Liveness Probe - Check application health
 - Readiness Probe - Check readiness to process requests


Liveness & Readiness

- Liveness Probe
 - Restarting containers if liveness probes fail
- Readiness Probe
 - Removing from service endpoint if readiness probe fails
- Probe methods
 - HTTP endpoint
 - TCP socket endpoint
 - Unix command's return value

```
apiVersion: v1
kind: Pod
metadata:
  name: app-with-readiness-check
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    livenessProbe:
      httpGet:
        path: /q/health
        port: 8080
      initialDelaySeconds: 30
    readinessProbe:
      exec:
        command: [ "stat", "/var/run/random-generator-ready" ]
```

Managed Lifecycle

How should applications react to lifecycle events?




Lifecycle Events

- SIGTERM
 - Initial event issued when a container is going to shutdown
 - Application should listen to this event to cleanup properly and then exit
- SIGKILL
 - Final signal sent after a grace period
 - Can't be caught by application
 - `terminationGracePeriodSeconds`
 - Time to wait after SIGTERM (default: 30s)

Lifecycle Hooks

- **postStart**
 - Called after container is created
 - Runs in parallel to the main container
 - Keeps Pod in status *Pending* until exited successfully
 - `exec` or `httpGet` handler types (like *Health Probe*)
- **preStop**
 - Called before container is stopped
 - Same purpose & semantics as for SIGTERM

```
apiVersion: v1
kind: Pod
metadata:
  name: post-start-pre-stop-hook
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    lifecycle:
      postStart:
        exec:
          command:
            - sh
            - -c
            - sleep 30 && echo "Wake up!" > /tmp/postStart_done
      preStop:
        httpGet:
          port: 8080
          path: /shutdown
```




Structural Patterns

How can I structure & organize my container(s) within a Pod?

Init Container

How can we initialize our containerized applications?

- Init Containers
 - Part of a Pod
 - One shot actions before application starts
 - Needs to be idempotent
 - Has own resource requirements




```
apiVersion: v1
kind: Pod
....
spec:
  initContainers:
  - name: download
    image: axec1br/git
    command: [ "git", "clone", "https://github.com/myrepo", "/data" ]
    volumeMounts:
    - mountPath: /var/lib/data
      name: source
  containers:
  - name: run
    image: docker.io/centos/httpd
    volumeMounts:
    - mountPath: /var/www/html
      name: source
  volumes:
  - emptyDir: {}
    name: source
```

Sidecar Pattern

How do we enhance the functionality of an application without changing it?

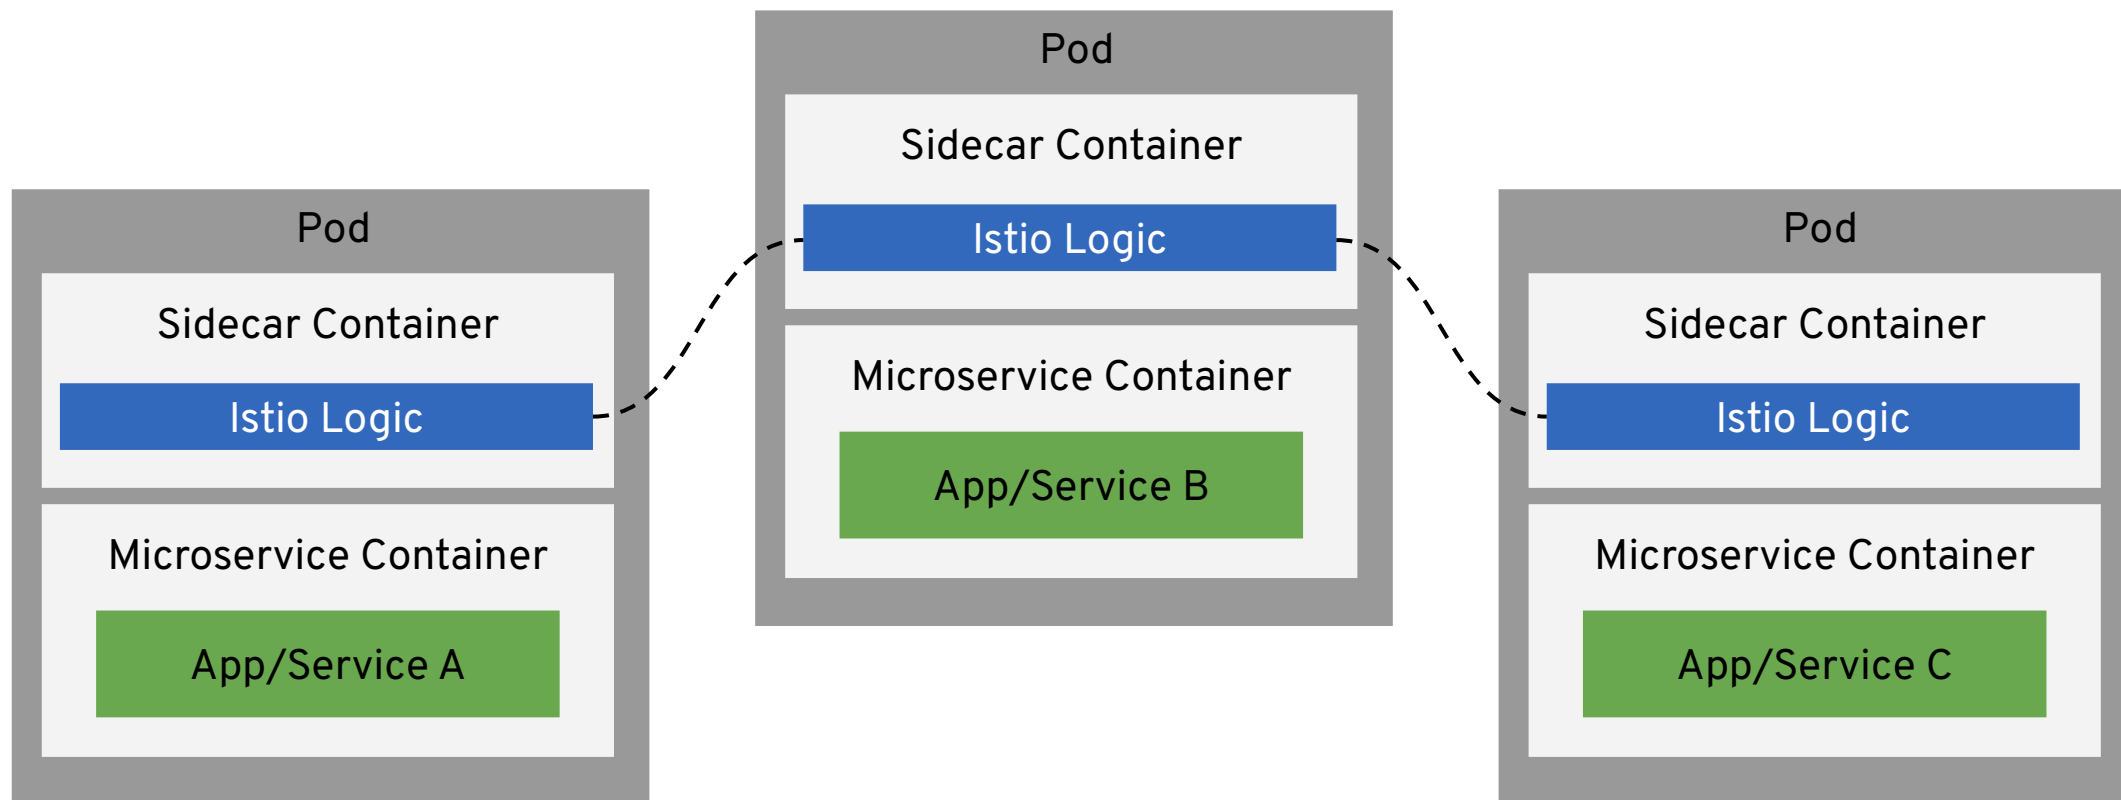
- Runtime collaboration of containers
- Connected via shared resources
 - Network
 - Volumes
- Similar what AOP is for programming
- Separation of concerns

Sidecar



```
apiVersion: v1
kind: Pod
....
spec:
  containers:
    - name: app
      image: docker.io/centos/httpd
      ports:
        - containerPort: 80
      volumeMounts:
        - mountPath: /var/www/html
          name: git
    - name: gitpoll
      image: axec1br/git
      volumeMounts:
        - mountPath: /var/lib/data
          name: git
      env:
        - name: GIT_REPO
          value: https://github.com/mdn/beginner-html-site-scripted
      command:
        - "sh"
        - "-c"
        - "git clone ${GIT_REPO} . && watch -n 600 git pull"
      workingDir: /var/lib/data
  volumes:
    - emptyDir: {}
      name: git
```

Service Mesh

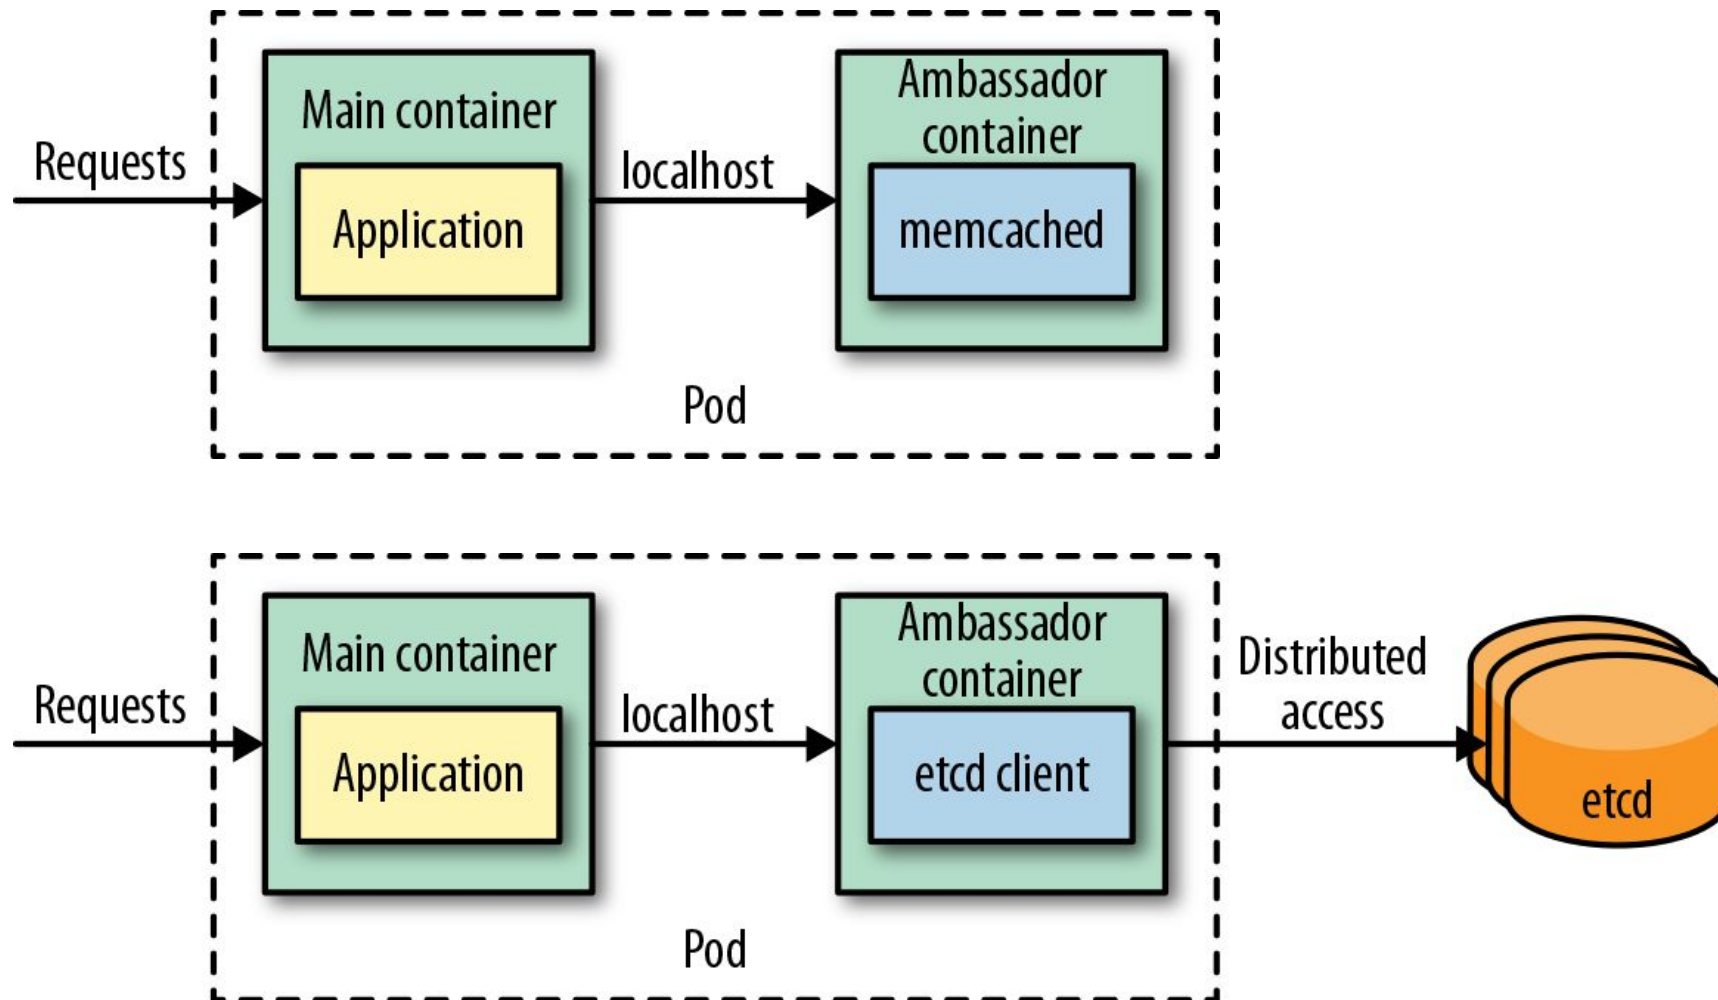


Ambassador Pattern

How to decouple a container's access to the outside world?

- Also known as **Proxy**
- Specialization of a Sidecar
- Examples for infrastructure services
 - Circuit breaker
 - Tracing

Ambassador




```
apiVersion: v1
kind: Pod
metadata:
  name: random-generator
  labels:
    app: random-generator
spec:
  containers:
    - image: k8spatterns/random-generator:1.0
      name: main
      env:
        - name: LOG_URL
          value: http://localhost:9009
      ports:
        - containerPort: 8080
          protocol: TCP
    - image: k8spatterns/random-generator-log-ambassador
      name: ambassador
```

Adapter Pattern

How to decouple access to a container from the outside world?

- Opposite of Ambassador
- Uniform access to an application
- Examples
 - Monitoring
 - Logging

Adapter



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: random-generator
spec:
  replicas: 1
  selector:
    matchLabels:
      app: random-generator
  template:
    metadata:
      labels:
        app: random-generator
    spec:
      containers:
        - image: k8spatterns/random-generator:1.0
          name: random-generator
          env:
            - name: LOG_FILE
              value: /logs/random.log
          ports:
            - containerPort: 8080
              protocol: TCP
          volumeMounts:
            - mountPath: /logs
              name: log-volume
```

```
  - image:
    k8spatterns/random-generator-exporter
      name: prometheus-adapter
      env:
        - name: LOG_FILE
          value: /logs/random.log
      ports:
        - containerPort: 9889
          protocol: TCP
      volumeMounts:
        - mountPath: /logs
          name: log-volume
      volumes:
        - name: log-volume
          emptyDir: {}
```



Configurational Patterns

How can applications be configured for different environments?

EnvVar Configuration

- Universally applicable
- Recommended by the *Twelve Factor App* manifesto
- **The Good**
 - Universal - every language understands them
- **The Bad**
 - Can only be set during startup of an application
 - Tedious if lots of them
 - Hard to know where they come from

ConfigMap

- Key-Value Map
- Use in Pods as
 - Environment variables
 - Volumes with keys as file names and values as file content

```
kubectl create cm spring-boot-config \  
  --from-literal=JAVA_OPTIONS=-Djava.security.egd=file:/dev/urandom \  
  --from-file=application.properties
```

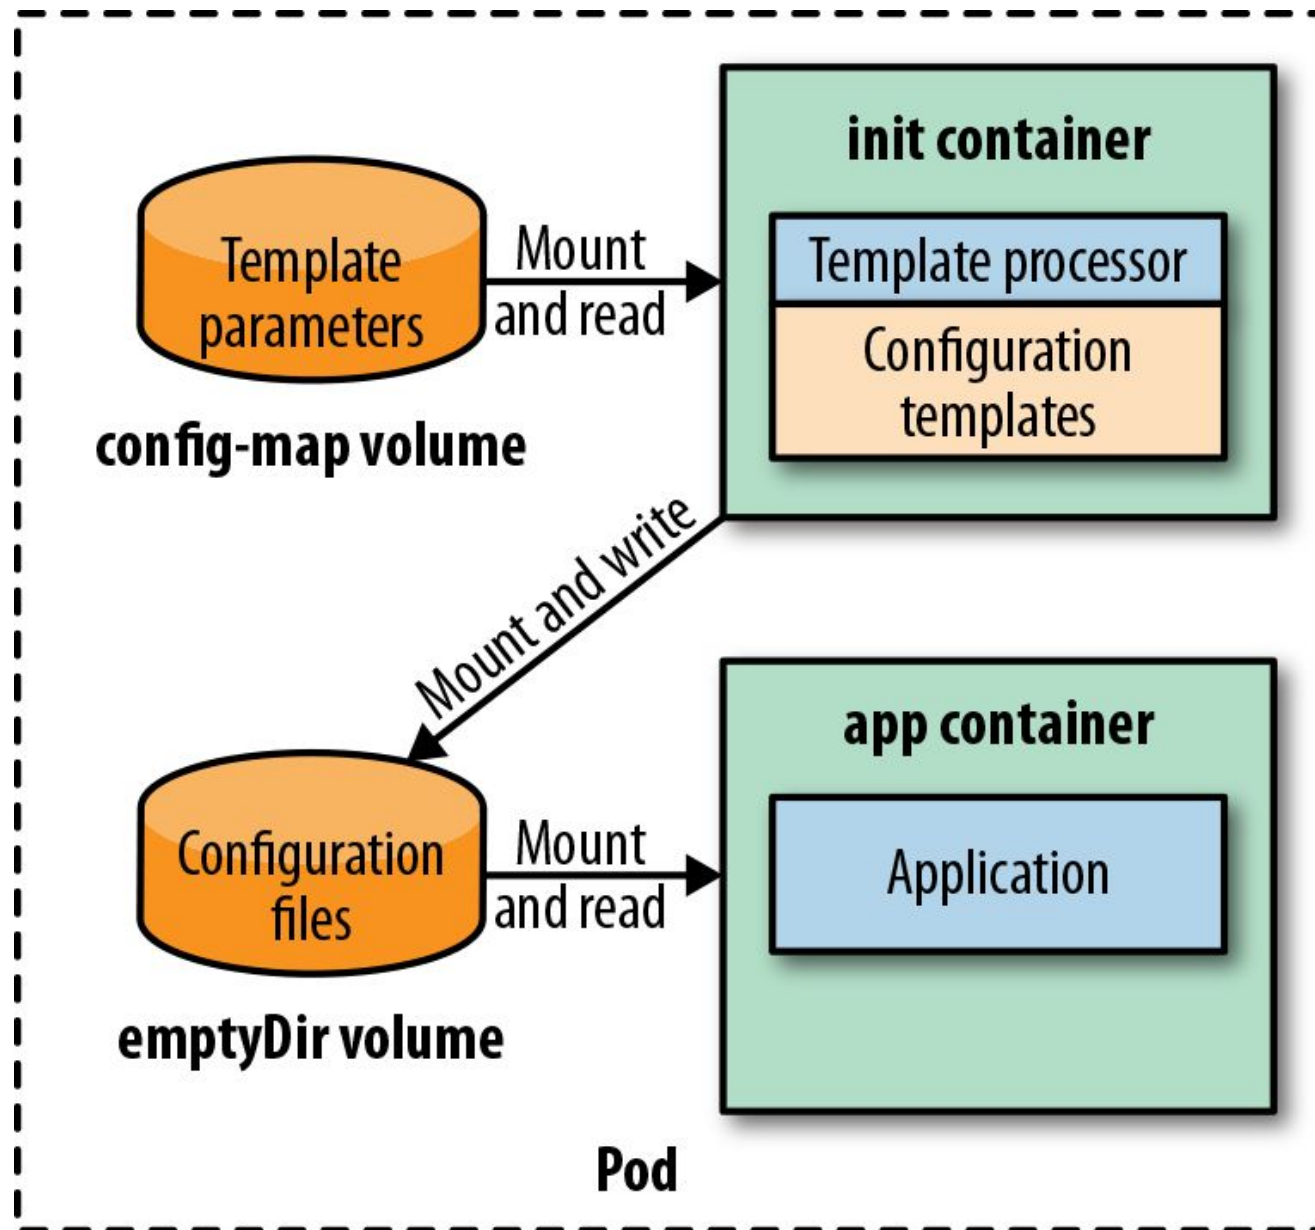
Secret

- Like **ConfigMap** but content Base64 encoded
- Secrets are ...
 - ... only distributed to nodes running Pods that need it
 - ... only stored in memory in a tmpfs and never written to physical storage
 - ... stored encrypted in the backend store (etcd)
- Access can be restricted with RBAC rules
- For high security requirements application based encryption is needed

Configuration Template

How to manage large and complex similar configuration data?

- **ConfigMap** not suitable for large configuration
- Managing similar configuration
- Ingredients
 - Init-container with template processor and templates
 - Parameters from a **ConfigMap** Volume
- Good for large, similar configuration sets per env
- Parameterization via **ConfigMaps** easy
- More complex




```
apiVersion: v1
kind: Deployment
metadata:
  name: wildfly-cm-template
spec:
  replicas: 1
  template:
    spec:
      initContainers:
        - image: k8spatterns/config-init
          name: init
          volumeMounts:
            - mountPath: "/params"
              name: wildfly-parameters
            - mountPath: "/out"
              name: wildfly-config
```

```
containers:
- image: jboss/wildfly:10.1.0.Final
  name: server
  volumeMounts:
    - mountPath: "/config"
      name: wildfly-config
volumes:
- name: wildfly-parameters
  configMap:
    name: wildfly-params-cm
- name: wildfly-config
  emptyDir: {}
```



Behavioral Patterns

How can my application
communicate/interact with
the underlying platform?

Self Awareness

How to gather information about a running Pod from within the Pod?

- Runtime data via Downward API
 - Pod name, ip address, hostname of underlying host
 - Specified resource requests & limits
 - Annotations & labels
- Application doesn't need to use a client
- Application can remain Kubernetes agnostic

```
apiVersion: v1
kind: Pod
metadata:
  name: random-generator
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    env:
    - name: POD_IP
      valueFrom:
        fieldRef:
        fieldPath: status.podIP
    - name: MEMORY_LIMIT
      valueFrom:
        resourceFieldRef:
        container: random-generator
```

```
apiVersion: v1
kind: Pod
metadata:
  name: random-generator
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    volumeMounts:
    - name: pod-info
      mountPath: /pod-info
    volumes:
    - name: pod-info
      downwardAPI:
        items:
        - path: labels
          fieldRef:
            fieldPath: metadata.labels
        - path: annotations
          fieldRef:
            fieldPath: metadata.annotations
```



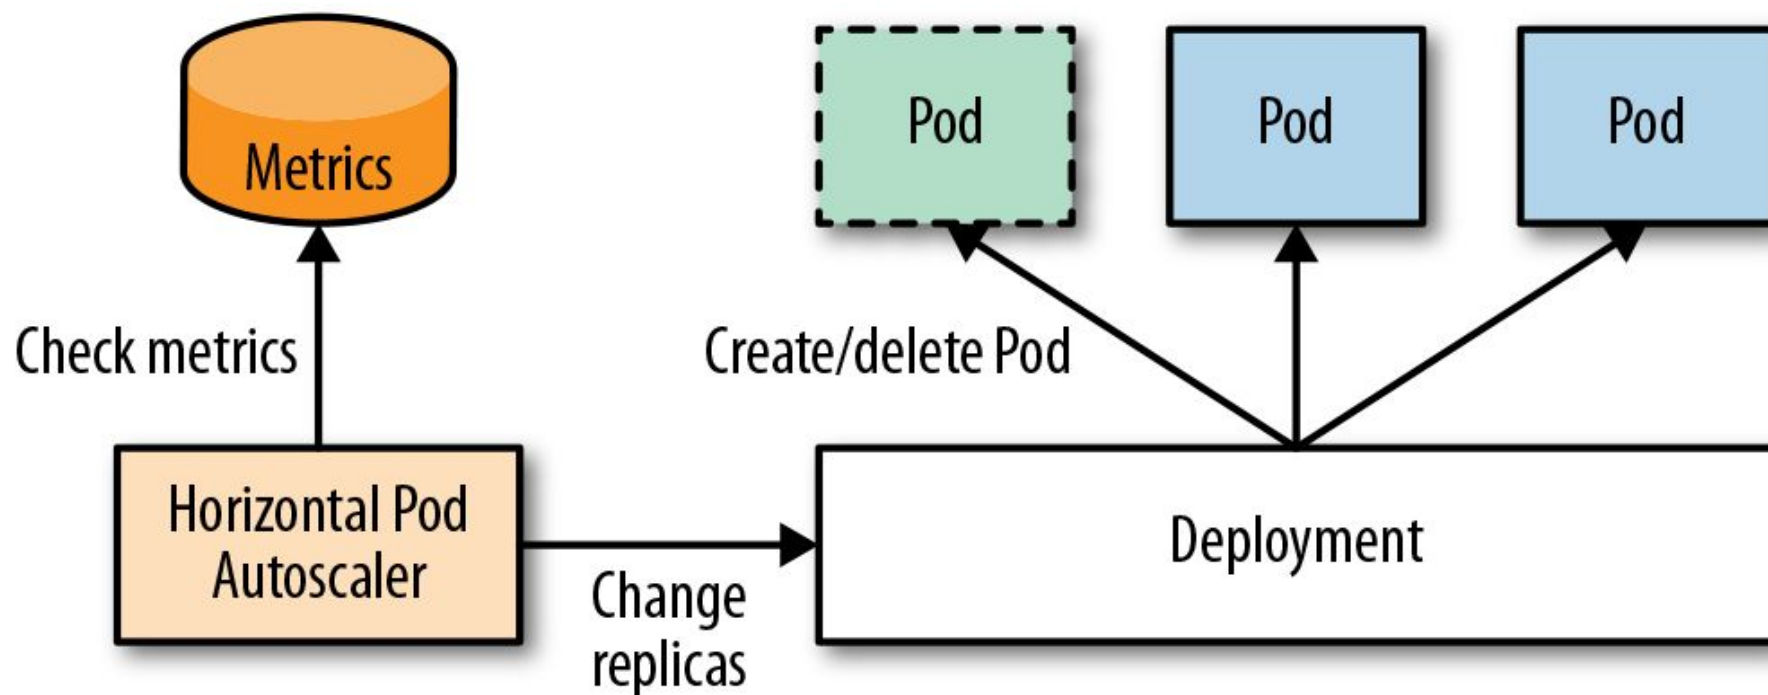
Advanced Patterns

Elastic Scale

How to automatically react to dynamically changing resource requirements?

- **Horizontal:** Changing replicas of a Pod
- **Vertical:** Changing resource constraints of containers in a single Pod
- **Cluster:** Adding new nodes to a cluster
- **Manual:** Changing scale parameters manually, imperatively, or declaratively
- **Automatic:** Change scaling parameters based on observed metrics

Horizontal Pod Autoscaler (HPA)



```
kubectl autoscale deployment random-generator --cpu-percent=50 --min=1 --max=5
```

HorizontalPodAutoscaler

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: random-generator
spec:
  minReplicas: 1
  maxReplicas: 5
  scaleTargetRef:
    apiVersion: extensions/v1beta1
    kind: Deployment
    name: random-generator
  metrics:
  - resource:
      name: cpu
      target:
        averageUtilization: 50
        type: Utilization
    type: Resource
```




Kubernetes **E**vent **D**riven **A**utoscaling

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: random-generator-so
spec:
  minReplicaCount: 1
  maxReplicaCount: 5
  scaleTargetRef:
    name: random-generator
  triggers:
  - type: kafka
    metadata:
      bootstrapServers: kafka.svc:9092
      consumerGroup: my-group
      topic: my-topic
      lagThreshold: '5'
```

Metrics & Challenges

- Metrics

- **Standard Metrics** - CPU & Memory Pod data obtained from Kubernetes metrics server
- **Custom Metrics** - Metrics delivered via an aggregated API server at the `custom.metrics.k8s.io` API path
- **External Metrics** - Metrics obtained from outside the cluster

- Challenges

- **Metric Selection** - Correlation between metric value and replica counts
- **Preventing Thrashing** - Windowing to avoid scaling on temporary spikes
- **Delayed Reaction** - Delay between cause and scaling reaction

A vertical red bar on the left side of the slide contains various white and dark red abstract icons representing technology, such as a cloud with a keyhole, a database cylinder, a server rack, a computer monitor, and various arrows and geometric shapes.

Thank you!

<https://k8spatterns.io/>

<https://github.com/k8spatterns/examples>



twitter.com/ro14nd



facebook.com/redhatinc



youtube.com/user/RedHatVideos

